

Aryan Sahni

 C0okiegranny221 |  Aryan Sahni |  Research |  aryn.sahni22@gmail.com |  +91 9958216312

SUMMARY

Research-driven Computer Science undergraduate at VIT Chennai with expertise in Artificial Intelligence, Machine Learning, Systems Programming, and Full-Stack Development. AI Intern with hands-on experience fine-tuning Veena TTS using the Indigo call data and designing a low-latency streaming architecture for a production-ready TTS service. Experienced in building full-stack AI platforms, Graph-RAG systems, and LLM-powered applications using TensorFlow, ONNX Runtime, LangChain, and modern AI frameworks. Published author at ImageCLEF Medical AI and multimodal image understanding, focused on developing scalable, data-driven solutions that bridge cutting-edge research with real-world applications.

EDUCATION

Vellore Institute of Technology, Chennai
B.Tech in Computer Science and Engineering

2023–2027
CGPA: 9.07/10

WORK EXPERIENCE

Data Science & AI Research Intern — AIONOS

[\[AIONOS\]](#)

- Engineered a **Chrome extension** for real-time **phishing and scam detection** leveraging the **Gmail API**, **Flask microservices**, and DOM event tracking via **MutationObserver**.
- Integrated **AWS Bedrock Foundation Models** for automated **LLM-powered threat classification**, contextual **email intent analysis**, and intelligent **response generation**.
- Implemented asynchronous frontend-backend communication pipelines and dynamic content parsing for low-latency inbox monitoring.

PROJECTS

GMAIL Bait Mail

[\[GitHub\]](#)

- Engineered a **Chrome extension** for real-time **phishing and scam detection** leveraging the **Gmail API**, **Flask microservices**, and DOM event tracking via **MutationObserver**.
- Integrated **AWS Bedrock Foundation Models** for automated **LLM-powered threat classification**, contextual **email intent analysis**, and intelligent **response generation**.
- Implemented asynchronous frontend-backend communication pipelines and dynamic content parsing for low-latency inbox monitoring.

Graph-RAG Knowledge Retrieval System

[\[GitHub\]](#)

- Architected a **Graph-RAG framework** using **LangChain**, **Ollama**, and **Zephyr LLMs** for scalable **multi-hop semantic retrieval**.
- Integrated **knowledge graph traversal** with **Retrieval-Augmented Generation (RAG)** pipelines to enhance contextual reasoning, entity linking, and response grounding.
- Developed custom evaluation modules for **faithfulness scoring**, **hallucination detection**, **graph coverage analysis**, and **LLM-as-a-judge benchmarking**.
- Optimized containerized deployment workflows and resolved **Docker–Ollama networking** bottlenecks for efficient local inference orchestration.

SKILLS

Programming Languages: Python, Java, C/C++, JavaScript, SQL, Bash

Tools & Platforms: Git/GitHub, Linux, Docker, ONNX Runtime

Frameworks & Libraries: TensorFlow, PyTorch, LangChain, Hugging Face Transformers, Pandas, NumPy, Scikit-learn

AI/ML Expertise: Generative AI, Large Language Models (LLMs), Speech AI & TTS, Deep Learning, NLP, Retrieval-Augmented Generation (RAG), Graph-RAG, Reinforcement Learning, Model Fine-Tuning, Model Evaluation

PUBLICATIONS

Sahni, Aryan, Rachit Gupta, et al. (2025). “Evaluating deep CNNs for multi-label concept detection in ROCov2 radiology image dataset by team LekshmiscscopeVIT”. In: *CLEF2025 Working Notes, CEUR Workshop Proceedings, CEUR-WS.org, Madrid, Spain*. URL: https://ceur-ws.org/Vol-4038/paper_203.pdf.

Sahni, Aryan, Krishnesh Mathur, et al. (2026). “An Optimal Evaluation of Entanglement Depth in Variational Quantum Classifiers for Efficient Deployment of Machine Learning Applications”. In: *2026 International Conference on Electronic Systems and Intelligent Computing (ICESIC)*, pp. 320–325. DOI: [10.1109/ICESIC67389.2026.11496417](https://doi.org/10.1109/ICESIC67389.2026.11496417).